

Lecture 9. The Shortest Path Problem

SIT221 Data Structures and Algorithms

Weighted Graph: Notation

- A graph is a pair $G = (V, E)$, where
 - V is a set of nodes, called **vertices**.
 - E is a collection of pairs of vertices, called **edges**.
 - We denote by $n = |V|$ the number of vertices and by $m = |E|$ the number of edges.
- Often there are **weights/costs** assigned to the edges: $c: E \rightarrow R$
- A graph is **weighted** if each edge is given a numerical weight.
Edge weights may represent, distances, costs, etc.

For example, in a flight route graph, the weight of an edge represents the distance in kilometers between the endpoint airports.

Shortest Path: Problem Formulation

Given a weighted directed graph $G = (V, E)$, and two vertices s and v , we want to find a path of minimum total cost among all possible paths between s and v .

Given a path $p = (e_1, e_2, \dots, e_k)$ consisting of k edges, the cost of the path is the sum of the weights of its edges:

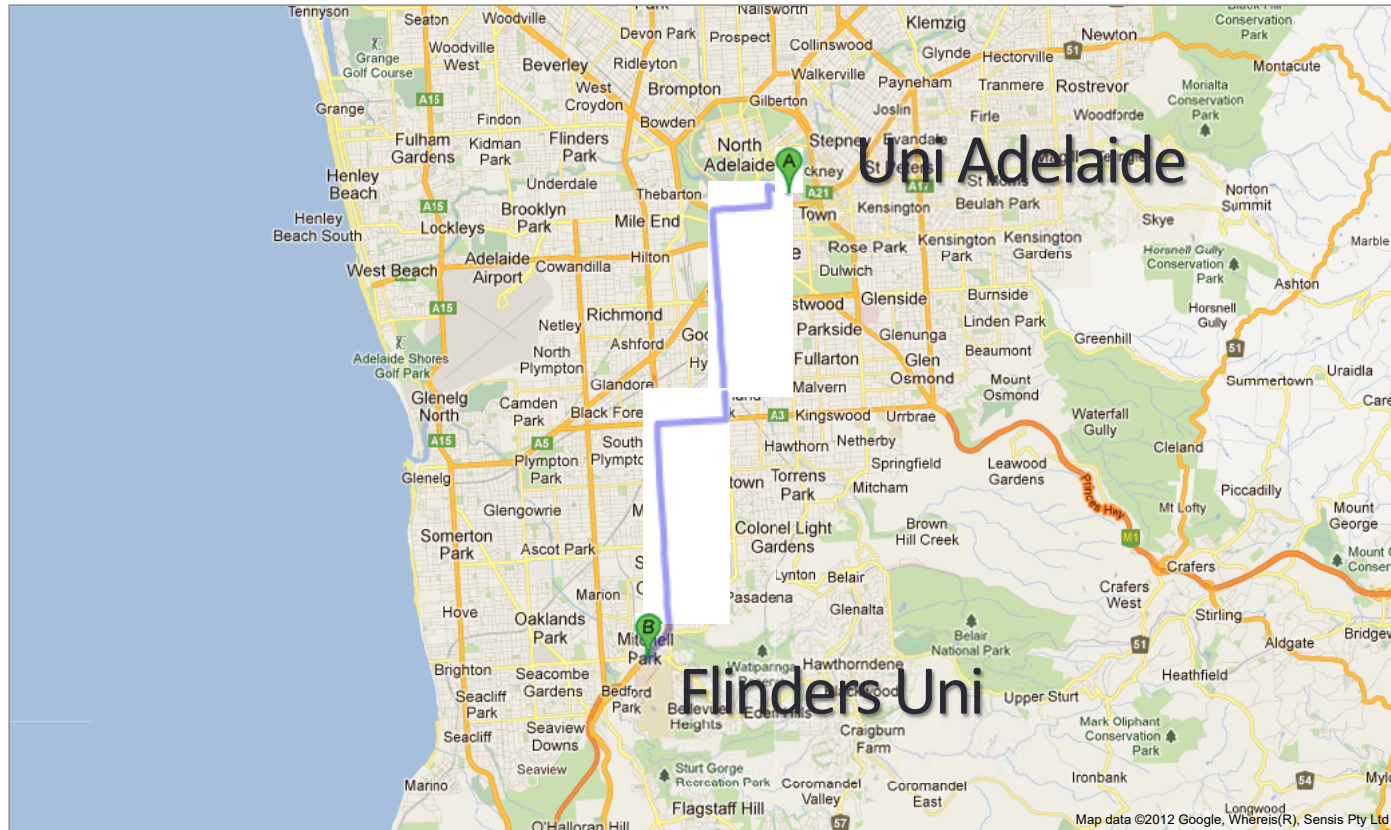
$$c(p) = \sum_{i=1}^k c(e_i).$$

Single-Source Shortest Path Problem:

Compute for a given node s of V a shortest path to any other node in V (if it exists).

We assume that edge weights are non-negative.

Shortest Path: Applications



Computation of shortest path is one of the classical problems.

- Route planning / driving directions
- Internet packet routing
- Flight reservations

Shortest Path: Dijkstra's Algorithm


- Remember BFS for computing all shortest paths in an unweighted graph.
- In iteration i , we computed all shortest paths having i edges.
- Dijkstra's algorithm obtains in iteration i a shortest path to the node of the i^{th} smallest distance from s .
- Dijkstra's algorithm computes the distances of all the vertices from a given start vertex s , assuming that:
 - the graph is connected
 - the edge weights are **nonnegative**

Dijkstra's Algorithm: Idea

- We grow a “**cloud**” of vertices, beginning with s and eventually covering all the vertices.
- We store with each vertex v a label $d(v)$ representing the distance of v from s in the subgraph consisting of the cloud and its adjacent vertices.
- At each step
 - We add to the cloud the vertex u outside the cloud with the smallest distance label, $d(u)$.
 - We update the labels of the vertices adjacent to u .


Dijkstra's Algorithm: Idea

We call a node u **unscanned** if no shortest path from s to u has been found so far.




Dijkstra's Algorithm: Idea

Make unscanned node u **scanned** that would get the minimal tentative distance among all unscanned nodes.




Dijkstra's Algorithm: Idea

Make unscanned node u **scanned** that would get the minimal tentative distance among all unscanned nodes.




Dijkstra's Algorithm: Idea

Consider all edges leaving u and update distances using the **Relax Routine**.



Dijkstra's Algorithm: Edge Relaxation

We may to update a previous path from s to v if we find a shorter path.




Procedure $relax(e = (u, v) : Edge)$

if $d[u] + c(e) < d[v]$ **then** $d[v] := d[u] + c(e); \quad parent[v] := u$

Dijkstra's Algorithm: Edge Relaxation

We may to update a previous path from s to v if we find a shorter path.



Procedure $relax(e = (u, v) : Edge)$

if $d[u] + c(e) < d[v]$ **then** $d[v] := d[u] + c(e); \quad parent[v] := u$

Dijkstra's Algorithm: Sketch

Dijkstra's Algorithm

declare all nodes unscanned


while there is an unscanned node with tentative distance $< +\infty$ **do**

$u :=$ the unscanned node with minimal tentative distance


 relax all edges (u, v) out of u and declare u scanned

Dijkstra's Algorithm: Solution Construction


1)




2)



3)




4)




Dijkstra's Algorithm: Solution Construction


5)





6)






Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example




-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached



Dijkstra's Algorithm: Example



-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached


Dijkstra's Algorithm: Example



-  Nodes for which a shortest path has been computed
-  Nodes that have already been reached

Dijkstra's Algorithm: Correctness

- Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.
- Suppose it did not find all shortest distances. Let F be the first wrong vertex the algorithm processed.
- When the previous node, D , on the true shortest path was considered, its distance was correct.
- But the edge (D, F) was relaxed at that time!
- Thus, so long as $d(F) > d(D)$, F 's distance cannot be wrong. That is, there is no wrong vertex.



Dijkstra's Algorithm: Correctness

Theorem: Dijkstra's algorithm solves the single-source shortest path problem for graphs with nonnegative edge costs.

Proof:

We show two steps:

- All nodes reachable from s are scanned after termination.
- When a node v becomes scanned then the shortest path from s to v is obtained.

Dijkstra's Algorithm: Correctness

Claim: All nodes reachable from s are scanned after termination.

Proof (by contradiction):

- Assume that there is a node v reachable from s , but never scanned.
- Consider a shortest path $p = (s = v_1, v_2, \dots, v_k = v)$ from s to v .
- Let $i > 1$ be minimal such that v_i is unscanned.
- Implies node v_{i-1} has been scanned.
- When v_{i-1} is scanned $d(v_i)$ is set to $d(v_{i-1}) + c(v_{i-1}, v_i) < \infty$.
- Hence, v_i must be scanned as only nodes u with $d(u) = \infty$ stay unscanned. Contradiction to v_i is unscanned.



Dijkstra's Algorithm: Correctness

Claim: When a node v becomes scanned then the shortest path from s to v is obtained.

Proof (by contradiction):

- Denote by $\mu(v)$ the length of a shortest path from s to v .
- Consider the first point in time t when v has been scanned and $d(v) > \mu(v)$ holds.
- Consider a shortest path $p = (s = v_1, v_2, \dots, v_k = v)$ from s to v .
- Let $i > 1$ be minimal such that v_i has not been scanned before time t .



Dijkstra's Algorithm: Correctness


Proof (continued):

- Node v_{i-1} was scanned before time t which implies $\mu(v_{i-1}) = d(v_{i-1})$.
- When v_{i-1} is scanned $d(v_i)$ is set to $d(v_{i-1}) + c(v_{i-1}, v_i) = \mu(v_{i-1}) + c(v_{i-1}, v_i)$.
- We have $d(v_i) = \mu(v_i) \leq \mu(v_k) < d(v_k)$ and hence v_i is scanned instead of v_k , **a contradiction**.



Dijkstra's Algorithm: Negative Weights

- Dijkstra's algorithm is based on the greedy method. It adds vertices by increasing distance.
- If a node with a negative incident edge were to be added late to the cloud, it could mess up distances for vertices already in the cloud.



C 's true distance is 1, but it is already in the cloud with $d(C) = 5$.

Dijkstra's Algorithm: Implementation

Store all unscanned reached nodes in an addressable priority queue Q (using tentative distances as key values)

- A priority queue stores the vertices outside the cloud
 - Key: distance
 - Element: vertex
- Locator-based methods
 - $\text{Insert}(k, e)$ returns a locator
 - $\text{DecreaseKey}(l, k)$ changes the key of an item
- We store two labels with each vertex:
 - Distance ($d(v)$ label)
 - Locator l in priority queue

Dijkstra's Algorithm: Pseudocode

Function *Dijkstra*($s : \text{NodeId}$) : $\text{NodeArray} \times \text{NodeArray}$

$d = \langle \infty, \dots, \infty \rangle : \text{NodeArray}$ **of** $\mathbb{R} \cup \{\infty\}$

$\text{parent} = \langle \perp, \dots, \perp \rangle : \text{NodeArray}$ **of** NodeId

$\text{parent}[s] := s$

$Q : \text{NodePQ}$

$d[s] := 0$; $Q.\text{insert}(s)$

while $Q \neq \emptyset$ **do**

$u := Q.\text{deleteMin}$

foreach $\text{edge } e = (u, v) \in E$ **do**

if $d[u] + c(e) < d[v]$ **then**

$d[v] := d[u] + c(e)$

$\text{parent}[v] := u$

if $v \in Q$ **then** $Q.\text{decreaseKey}(v)$

else $Q.\text{insert}(v)$

return (d, parent)



// returns (d, parent)
// tentative distance from root

// self-loop signals root
// unscanned reached nodes

// we have $d[u] = \mu(u)$

// relax

// update tree



Dijkstra's Algorithm: Runtime Complexity

- Initialization (arrays, priority queue) takes time $O(n)$.
- Every reachable node is inserted and removed once from Q .
- At most n DeleteMin and insert operations.
- Each node is scanned at most once and each edge is relaxed at most once.
- Implies at most m DecreaseKey operations.

Total runtime

$$T_{\text{Dijkstra}} = O\left(m \cdot T_{\text{decreaseKey}}(n) + n \cdot (T_{\text{deleteMin}}(n) + T_{\text{insert}}(n))\right)$$

Dijkstra's Algorithm: Runtime Complexity

Runtime depends on implementation of priority queue.

Original (Dijkstra 1959):

- Maintain the number of reached unscanned nodes.
- An array d storing the distances and an array storing for each node whether it is reached or unscanned.
- Insert and DecreaseKey take time $O(1)$
- DeleteMin takes time $O(n)$
- Total Runtime: $O(m + n^2)$

Improvements:

- Binary Heaps: $O((m + n) \log n)$
- Fibonacci Heaps: $O(m + n \log n)$

Shortest Path: Properties

Property 1: There is a tree of shortest paths from a start vertex to all the other vertices.

Property 2: A subpath of a shortest path is itself a shortest path.

Proof (by contradiction):

Assume that the path p is a shortest path from s to v .

Assume that a subpath from a to b is not a shortest path from a to b




This implies that there is a shorter path from a to b

We can use this path to obtain a shorter path from s to v .

Contradiction to p is shortest path from s to v .



Shortest Path: Negative Cycles




If a path from s to v contains a negative cycles then a shortest path does not exist (is not defined).

Bellman-Ford Algorithm


- Dijkstra's algorithm works for acyclic graphs and for non-negative edge costs.
- Bellman-Ford algorithm solves the problem for arbitrary edge costs.
- It uses $n - 1$ rounds and relaxes in each round all edges.
- This works as simple paths have at most $n - 1$ edges.
- After the relaxations are complete, we have all shortest paths to nodes with non-negative cycles.
- We still need to identify the nodes that can be reached by using negative cycles.

Bellman-Ford Algorithm: Solution Construction


1)




2)



3)



4)



Nodes are labeled with their $d(v)$ values

Bellman-Ford Algorithm: Negative Weights

- Assume that there is an edge $e = (u, v)$ that allows to improve $d(v)$ after the relaxations are complete.
- Then the node v is reachable by using a negative cycle.
- Furthermore, all nodes reachable from v can also be reached by using a negative cycle.
- We set $d(v) = -\infty$ for these nodes v .
- We use postprocessing and the routine **Infect** to find nodes reachable by negative cycles.

Bellman-Ford Algorithm: Pseudocode

Function *BellmanFord*($s : \text{NodeId}$) : $\text{NodeArray} \times \text{NodeArray}$

$d = \langle \infty, \dots, \infty \rangle : \text{NodeArray}$ **of** $\mathbb{R} \cup \{-\infty, \infty\}$

// distance from root

$parent = \langle \perp, \dots, \perp \rangle : \text{NodeArray}$ **of** NodeId

$d[s] := 0$; $parent[s] := s$

// self-loop signals root

for $i := 1$ **to** $n - 1$ **do**

forall $e \in E$ **do** *relax*(e)

// round i

forall $e = (u, v) \in E$ **do**

// postprocessing

if $d[u] + c(e) < d[v]$ **then** *infect*(v)

return ($d, parent$)

Procedure *infect*(v)


if $d[v] > -\infty$ **then**

$d[v] := -\infty$

foreach $(v, w) \in E$ **do** *infect*(w)

Running time: $O(nm)$.

Shortest Path: Exercise



Other references and things to do

- Read chapter 14.6 in Data Structures and Algorithms in Java. Michael T. Goodrich, Irvine Roberto Tamassia, and Michael H. Goldwasser, 2014.