

Lecture 7. Heap Sort. AVL-Trees. Graphs and their Representation.

SIT221 Data Structures and Algorithms

Sorting with Priority Queues

- We use a Priority Queue
 - Insert the elements with a series of **Insert** operations.
 - Remove the elements in sorted order with a series of **DeleteMin** operations.
- The running time depends on the priority queue implementation:
 - Unsorted sequence gives the Selection Sort: $O(n^2)$ time
 - Sorted sequence gives Insertion Sort: $O(n^2)$ time
- Can we do better?

Sorting with Priority Queues

Algorithm *PriorityQueueSort*(S, C)

Input: sequence S , comparator C for the elements of S

Output: sequence S sorted in increasing order according to C

// Build priority queue P applying comparator C

while (**not** S .isEmpty()) **do**

 Element $e = S$.First();

P .Insert(e);

S .Remove(e);

// Build back the (sorted) sequence S

while (**not** P .isEmpty()) **do**

 Element $e = P$.DeleteMin();

S .AddLast(e);

Sorting with a Minimum Binary Heap: Heap Sort

Want to have a sorting algorithm based on heaps that runs in time $O(n \log n)$.

Idea:


- Build (the bottom-up strategy) the heap for n elements in time $O(n)$.
- Pick in each step the minimum element and delete it in time $O(\log n)$.
- Iterate until heap is empty.
- The space used is $O(n)$.

In total, n iterations implies the total runtime of $O(n \log n)$

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 1. Bottom-up heap construction



value:

8	10	4	9	14	7	1
---	----	---	---	----	---	---


indices:

1	2	3	4	5	6	7
---	---	---	---	---	---	---

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 1. Bottom-up heap construction



value:

8	10	4	9	14	7	1
---	----	---	---	----	---	---


indices:

1	2	3	4	5	6	7
---	---	---	---	---	---	---

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 2. Iterative Deletion: Delete 1



value:

4	9	7	10	14	8	
---	---	---	----	----	---	--


Sorted array:

1						
---	--	--	--	--	--	--

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 2. Iterative Deletion: Delete 4



value:

7	9	8	10	14		
---	---	---	----	----	--	--


Sorted array:

1	4					
---	---	--	--	--	--	--

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 2. Iterative Deletion: Delete 7



value:

8	9	14	10			
---	---	----	----	--	--	--


Sorted array:

1	4	7				
---	---	---	--	--	--	--

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 2. Iterative Deletion: Delete 8



value:

9	10	14				
---	----	----	--	--	--	--


Sorted array:

1	4	7	8			
---	---	---	---	--	--	--

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 2. Iterative Deletion: Delete 9



value:

10	14					
----	----	--	--	--	--	--

Sorted array:

1	4	7	8	9		
---	---	---	---	---	--	--

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 2. Iterative Deletion: Delete 10

14

value:

14						
----	--	--	--	--	--	--

Sorted array:

1	4	7	8	9	10	
---	---	---	---	---	----	--

Heap Sort: Example

Sort the sequence [8,10,4,9,14,7,1]

Step 2. Iterative Deletion: Delete 14

value:

--	--	--	--	--	--	--

Sorted array:

1	4	7	8	9	10	14
---	---	---	---	---	----	----

Heap Sort: Properties and Complexity

- Heapsort is in-place, but is **not** a stable sort.
- Requires only a constant amount of auxiliary space, i.e. less than the Merge Sort needs.
- Slower in practice on most machines than a well-implemented Quick Sort, it has the advantage of a more favourable worst-case $O(n \log n)$ runtime.

- Worst case: $T(n) = O(n \log n)$ comparisons
- Best case: $T(n) = O(n \log n)$ comparisons
- Average case: $T(n) = O(n \log n)$ comparisons
- Worst-case space complexity $O(1)$ auxiliary

Short Summary of Sorting Algorithms

Algorithm	Time	Notes
Selection Sort	$O(n^2)$	slow in-place for small data sets (< 1K)
Insertion Sort	$O(n^2)$	slow in-place for small data sets (< 1K)
Heap Sort	$O(n \log n)$	fast in-place for large data sets (1K — 1M)
Merge Sort	$O(n \log n)$	fast sequential data access for huge data sets (> 1M)

Runtimes for Binary Search Tree

Find, insert, and remove:

- Worst case: $\theta(n)$
- Best case: $\theta(\log n)$ and $\theta(1)$ for Find(k)
- Average case: $\theta(\log n)$


Observation: Binary search trees can get imbalanced when applying insert and/or remove operations.

Aim: Time $O(\log n)$ in the worst case for all operations

Idea: Whenever a subtree rooted at a node v gets imbalanced, apply operations that balance it out in time $O(\log n)$.

AVL-Tree: Definition

- An **AVL-Tree** is a binary search tree such that for every internal node v of T , the heights of the children of v can differ by at most 1.
- AVL-trees are balanced.
- The height of an AVL-Tree storing n keys is $O(\log n)$.



An example of an AVL-tree where the heights are shown next to the nodes.

AVL-Tree: Formal Definition


- Let $h(T)$ be the height of a tree T .
- Let v be a node in T , and T_l and T_r be the left and right subtree of v .
- We denote by $b(v) = h(T_l) - h(T_r)$ the balance degree of v .

Definition: A binary search tree T is called an AVL-tree if for each $v \in T$, $b(v) \in \{-1, 0, 1\}$ holds.

AVL-Tree: Insertion

- Insertion is as in a binary search tree.
- Always done by expanding an external node.

Example:




Insert 7

AVL-Tree: Insertion

- Insertion is as in a binary search tree
- Always done by expanding an external node.

Example:




Consider path from new leaf to the root
and check the balance values

AVL-Tree: Insertion

- Insertion is as in a binary search tree
- Always done by expanding an external node.


Example:



AVL-Tree: Insertion

- Insertion is as in a binary search tree
- Always done by expanding an external node.

Example:



Rotation establishes AVL-property again

AVL-Tree: Rebalancing after Insertion

- Inserting a new element z can violate the AVL-property.
- Consider path from the newly inserted leaf z to the root.
- Update the balance values.
- Repair AVL-property (if necessary).

AVL-Tree: Rebalancing after Insertion

- We insert new node z as for Binary Search Trees.
- $b(z) = 0$ holds after insertion.
- $b(v)$ might change by 1 for a node v on the path from z to the root.
- If $b(v) \notin \{-1, 0, 1\}$ then rebalance.

AVL-Tree: Rebalancing (Left Rotation)

Assume we have added z into the right subtree of node v .
Start examining for v , where v is the parent of z , and continue with the parent of v (if necessary).

Before insertion \rightarrow After Insertion:

- $b(v) = 1 \rightarrow b(v) = 0$
(height of tree rooted at v has not changed, stop rebalancing)
- $b(v) = 0 \rightarrow b(v) = -1$
(height of tree rooted at v has increased by 1, stop rebalancing only if v is root, otherwise examine parent of v)
- $b(v) = -1 \rightarrow b(v) = -2$
(AVL-property violated, carry out rotation)


AVL-Tree: Left Rotation

Assume node v and right child x of node v is on the path from z to the root.

w denotes the right child of x on the path

\Rightarrow Left rotation

New balance values: $b(x) = 0$ and $b(v) = 0$




AVL-Tree: Right Rotation

Assume node v and left child x of node v is on the path from z to the root.

w denotes the left child of x on the path


\Rightarrow Right rotation

New balance values: $b(x) = 0$ and $b(v) = 0$




AVL-Tree: Right-Left Rotation

w is left child of x on the path \Rightarrow Right-Left Rotation.



AVL-Tree: Right-Left Rotation

w is right child of x on the path \Rightarrow Left-Right Rotation.




AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6




AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6




AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6



AVL-Tree: Insertion


Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6



Balance OK


AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6




AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6




AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6



AVL-Tree: Insertion


Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6



Unbalanced
⇒ Left-right Rotation


AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6




AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6



AVL-Tree: Insertion

Example: Create AVL-Tree for sequence 4, 5, 7, 2, 1, 3, 6



AVL-Tree: Algorithm to perform Rotations

- **IF** (*tree is right heavy*)
 - **IF** (*tree's right subtree is left heavy*) perform Right-Left Rotation
 - **ELSE** perform Single Left Rotation
- **ELSE IF** (*tree is left heavy*)
 - **IF** (*tree's left subtree is right heavy*) perform Left-Right Rotation
 - **ELSE** perform Single Right Rotation

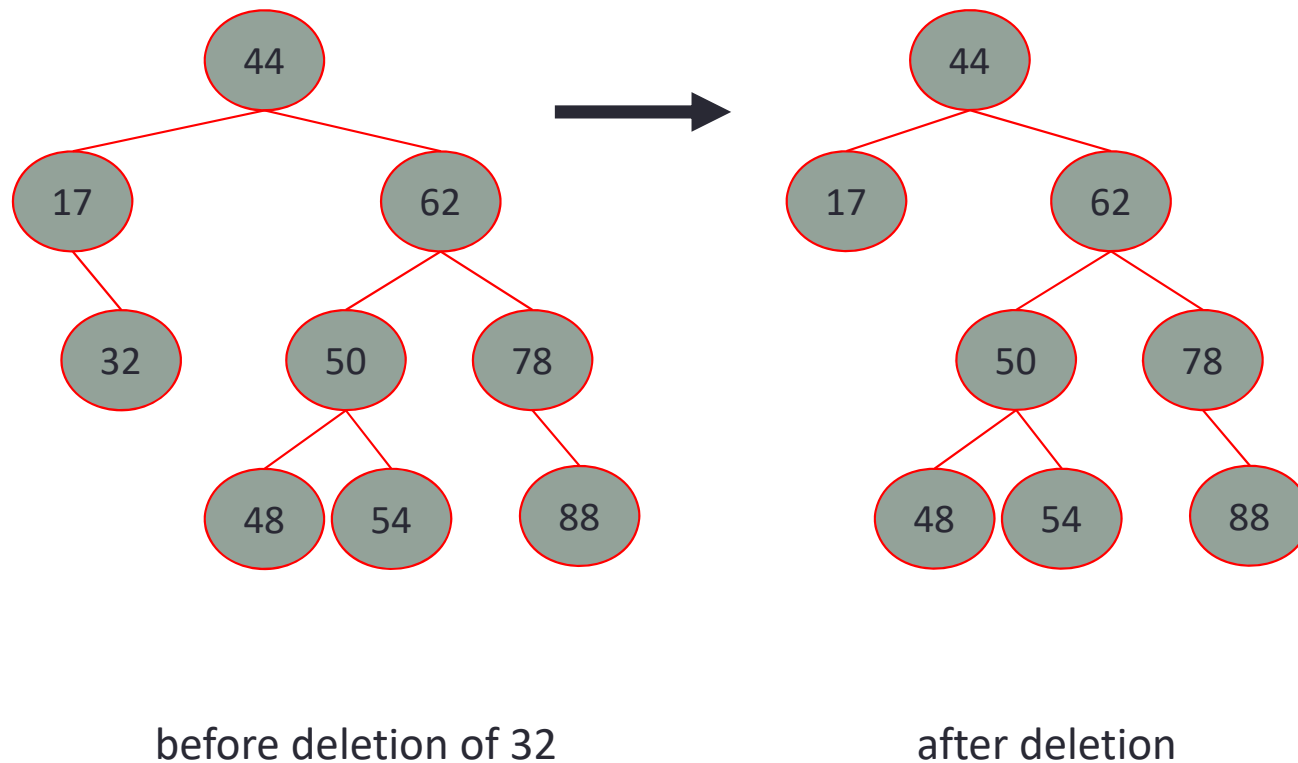
IMPORTANT: Maintain the Binary Search Tree Property

Let u , v , and w be three nodes such that u is in the left subtree of v and w is in the right subtree of v . Then we have

$$\text{key}(u) \leq \text{key}(v) \leq \text{key}(w)$$


AVL-Tree: Deletion

Removal begins as in a binary search tree, which means the node removed will become an empty external node. Its parent, v , may cause an imbalance.



AVL-Tree: Left Rotation after Deletion


Assume that the deleted node was in the left subtree of v and height of this tree has decrease by 1.



If B had height $h-1$ before deletion, the height of the subtree has decreased


AVL-Tree: Right Rotation after Deletion

Assume that the deleted node was in the right subtree of v and height of this tree has decrease by 1.



If B had height $h-1$ before deletion, the height of the subtree has decreased


AVL-Tree: Right-Left Rotation after Deletion



The height of the subtree has decreased

Either B or C might have height $h-1$

AVL-Tree: Left-Right Rotation after Deletion



The height of the subtree has decreased

Either B or C might have height $h-1$

Rebalancing after Deletion

- After having rebalanced for node v the height of the tree previously rooted at v might have decreased after deleting and rebalancing.
- If this is the case, old parent of v might be imbalanced.
- We might have to **continue rebalancing until the root has been reached**.

Running Times for AVL Trees

- A single restructure is $O(1)$
 - Using a linked-structure binary tree
- Find is $O(\log n)$
 - Height of tree is $O(\log n)$, no restructures needed
- Insert is $O(\log n)$
 - Initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$
- Remove is $O(\log n)$
 - Initial find is $O(\log n)$
 - Restructuring up the tree, maintaining heights is $O(\log n)$


Graph: Terminology and Representations




The metropolitan area of Milan, Italy at night. Astronaut photograph ISS026-E-28829, 2011. U.S. government image. NASA-JSC.

Graph: Terminology and Representations


- A graph is a pair $G = (V, E)$, where
 - V is a set of nodes, called **vertices**.
 - E is a collection of pairs of vertices, called **edges**.
 - We denote by $n = |V|$ the number of vertices and by $m = |E|$ the number of edges.
- Example:
 - A vertex represents an airport and stores the three-letter airport code.
 - An edge represents a flight route between two airports and stores the mileage of the route.



Graph: Edge Types



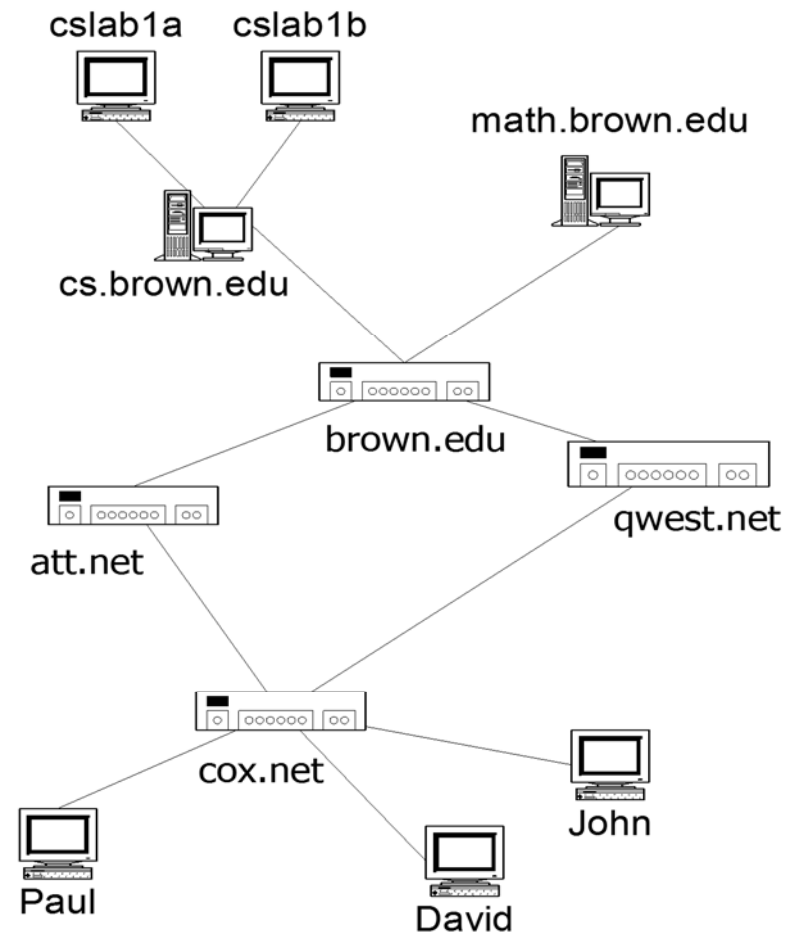
- **Directed edge**
 - ordered pair of vertices (u, v)
 - first vertex u is the origin
 - second vertex v is the destination
 - e.g., a flight
- **Directed graph**
 - all the edges are directed
 - e.g., route network



- **Undirected edge**
 - unordered pair of vertices (u, v)
 - e.g., a flight route
- **Undirected graph**
 - all the edges are undirected
 - e.g., flight network

Graph: Applications

- Electronic circuits
 - Printed circuit board
 - Integrated circuit
- Transportation networks
 - Highway network
 - Flight network
- Computer networks
 - Local area network
 - Internet
 - Web
- Databases
 - Entity-relationship diagram



Graph: Terminology

Adjacent vertices:
 u and v are adjacent


Edges **incident** on a vertex:
 a , d , and b are incident on v

Degree of a vertex:
 x has degree 5

Self-loop:
edge j is a self-loop


Parallel edges:
 h and i are parallel edges

End vertices (or **endpoints**)
of an edge: w and y are the
endpoints of edge f




Graph: Terminology

- Path
 - sequence of alternating vertices and edges
 - begins with a vertex
 - ends with a vertex
 - each edge is preceded and followed by its endpoints
- Simple path
 - path such that all its vertices and edges are distinct
- Examples
 - $P_1 = (v, b, x, h, z)$ is a simple path
 - $P_2 = (u, c, w, e, x, g, y, f, w, d, v)$ is a path that is not simple



Graph: Terminology

- Cycle
 - circular sequence of alternating vertices and edges
 - each edge is preceded and followed by its endpoints
- Simple cycle
 - cycle such that all its vertices and edges are distinct
- Examples
 - $C_1 = (v, b, x, g, y, f, w, c, u, a)$ is a simple cycle
 - $C_2 = (u, c, w, e, x, g, y, f, w, d, v, a)$ is a cycle that is not simple



Graph: Terminology

- The number of outgoing edges of a vertex v is called the **outdegree** of v :

$$\text{outdegree}(v) = |\{(v, u) \in E\}|$$

- The number of incoming edges of a vertex v is called the **indegree** of v :

$$\text{indegree}(v) = |\{(u, v) \in E\}|$$

Graph: Terminology

- A graph $G' = (V', E')$ is a subgraph of $G = (V, E)$ if

$$V' \subseteq V \text{ and } E' \subseteq E.$$

- Given a graph $G = (V, E)$ and a subset $V' \subseteq V$, the subgraph induced by V' is defined as

$$G' = (V', E \cap (V' \times V'))$$

Graph: Simple Graph Algorithm

Given a directed graph $G = (V, E)$. Is G acyclic?

Observation:

Node with outdegree zero can not appear in a cycle.

Idea for an algorithm:

- If there is a node v with outdegree zero, delete v (and the incoming edges) to obtain a graph G' ;
- G is acyclic if and only if G' is acyclic.


Graph: Simple Graph Algorithm

- If there is a node v of outdegree zero, delete v and its incoming edges to obtain a graph G' .
- Iterate the transformation.


Arrive at a graph G^*

- If G^* is the empty graph then G is acyclic;
- If G^* is not the empty graph, we can find a cycle in G^* that is also present in G .


Graph: Graph containing a cycle



Graph: Graph containing a cycle



Graph: Acyclic Graph



Empty Graph G^* implies that G is acyclic

Graph: Trees

An undirected graph is called a **tree** if there is **exactly one path between any pair of nodes**.

The following properties of an undirected graph G are equivalent:

1. G is a tree.
2. G is connected and has exactly $n - 1$ edges.
3. G is connected and contains no cycles.

Graph: Operations

We want efficiently support the following operations for graphs:

- **Accessing associated information**
(get the information stored at nodes and edges)
- **Navigation** (access the edges incident to a node)
- **Edge queries** (ask whether an edge is in the graph, query its reverse edge)
- **Construction, conversion and output**
(translate one graph representation into another)
- **Update** (Add and remove nodes and edges)

Graph: Representation

Simplest choice:

Unordered sequence of edges
(e.g. linked list of edges).

Good if you just want to output the edges of the graph.


Problem:

Most interesting operations take time $\Theta(m)$.


Graph: Adjacency Arrays

- Assume that the graph is static (i.e. it does not change).
- Then we can store the graph in an array.
- Store the outgoing neighbors of each node in a subarray and concatenate these subarrays into a single edge array E .
- Use an additional array V to store the starting positions of the subarrays.
- Memory consumption: $n + m + \Theta(1)$.

Graph: Adjacency Arrays



(Bi)-directed Graph



Adjacency Array

- For any node v , $V[v]$ is the index of the first outgoing edge of v .
- Add dummy entry $V[n + 1] = m + 1$.
- Outgoing edges of node v are accessible at $E[V[v]], \dots, E[V[v + 1] - 1]$

Graph: Adjacency Arrays


Are there better representations that allow to add or remove edges in constant time?

Two popular choices:


Adjacency Lists

Adjacency Matrices

Graph: Adjacency Arrays



(Bi)-directed Graph



Adjacency List

$$\begin{pmatrix} 0 & 1 & 0 & 1 \\ 1 & 0 & 1 & 1 \\ 0 & 1 & 0 & 1 \\ 1 & 1 & 1 & 0 \end{pmatrix}$$

Adjacency Matrix

Graph: Adjacency Arrays

Idea: Use for each node v a double-linked list that stores its outgoing neighbors (alternatively we can also use the incoming neighbors or lists for both).

Advantage:

- Insertion of edges goes in constant time.
- Well suited for sparse graphs (occur often in practice).

Graph: Adjacency Matrices

Idea: Represent a graph consisting of v nodes by an $v \times v$ matrix A . Set

$$A_{ij} = 1 \text{ if } (i, j) \in E$$
$$A_{ij} = 0 \text{ otherwise}$$

Insertion, removal, edge queries work in constant time.
 $O(n)$ to obtain an edge entering or leaving a node.

Disadvantage: Storage requirement n^2 even for sparse graphs.

Other references and things to do

- Have a look at the attached references in CloudDeakin.
- Read chapters 9.4.2, 11.3, and 14.2 in Data Structures and Algorithms in Java. Michael T. Goodrich, Irvine Roberto Tamassia, and Michael H. Goldwasser, 2014.